



FineWeb. decanting the web for the finest text data at scale

HuggingFaceFW / [blogpost-fineweb-v1](#)

1.38k



FineWeb

The finest collection of data the web has to offer



AUTHORS

Guilherme Penedo, Hynek Kydlíček, Loubna Ben Allal, Anton Lozhkov, Colin Raffel, Leandro Werra, Thomas Wolf

AFFILIATION

HuggingFace

PUBLISHED

May 31, 2024

The performance of a large language model (LLM) depends heavily on the quality and size of its pretraining dataset. However, the pretraining datasets for state-of-the-art open LLMs like Llama 3 [1] and Mixtral [2] are not publicly available and very little is known about how they were created.

Reading time: 45 min. For the best reading experience, we recommend not using a mobile phone.

Recently, we released [🍷 FineWeb](#), a new, large-scale (**15-trillion tokens, 44TB disk space**) dataset for LLM pretraining. FineWeb is derived from 96 [CommonCrawl](#) snapshots and produces **better-performing LLMs than other open pretraining datasets**. To bring more clarity in machine learning and advance the open understanding of how to train good quality large language models, we carefully documented and ablated all of the design choices used in FineWeb, including in-depth investigations of deduplication and filtering strategies. The present long form report is a deep dive in how to create a large and high-quality web-scale dataset for LLM pretraining. The dataset itself, [🍷 FineWeb](#), is available [here](#).

We are extremely thankful to the whole [distill.pub](#) team (Christopher Olah, Shan Carter, Ludwig Schubert in particular) for creating the template on which we based this blog post. Thanks also for inspiring us with exquisitely crafted articles and blog posts.